

1 The time hierarchy theorem

Theorem 1. *There's a universal $c > 1$ such that for any time-constructible $T(n) \geq n$ we have $\mathcal{DTIME}[T^c] \not\subseteq \mathcal{DTIME}[T]$.*

Proof. Let U be the universal machine, and let $c_U > 1$ be the constant such that if $M(x)$ runs in time T then $U(\langle M \rangle, x)$ runs in time at most T^{c_U} .

Defining a language $L \subseteq \{0,1\}^*$. For convenience, we define L by defining an indicator function $L: \{0,1\}^* \rightarrow \{0,1\}$ (i.e., x is in the language iff the indicator function outputs 1 on x). For any $x \in \{0,1\}^*$, if x isn't a valid encoding $\langle M \rangle$ of a TM, then $L(x) = 0$. Otherwise, consider the computation of $U(\langle M \rangle, \langle M \rangle)$ truncated after $T(|\langle M \rangle|)^{c_U}$ steps. If U halted and returned a bit b at that point (or before), then $L(x) = 1 - b$. Otherwise (if U didn't halt), then $L(x) = 0$.

A lower bound: $L \notin \mathcal{DTIME}[T]$. Assume towards a contradiction that a T -time M decides L . Then $M(\langle M \rangle)$ stops after at most $T(|\langle M \rangle|)$ steps and outputs some bit b , in which case $U(\langle M \rangle, \langle M \rangle)$ stops after at most $T(|\langle M \rangle|)^{c_U}$ steps and outputs b . Hence $L(\langle M \rangle) = 1 - b$. Then, we have

$$b = M(\langle M \rangle) = L(\langle M \rangle) = 1 - b,$$

a contradiction (where the middle inequality is since M decides L).

An upper bound: $L \in \mathcal{DTIME}[T^c]$. We design a three-tape M_L that decides L in time $\text{poly}(T)$ (and rely on the simulation of three-tape TMs by standard TMs).

If the input $x \in \{0,1\}^n$ isn't a valid encoding $\langle M \rangle$ of a TM, M_L rejects. Otherwise, M_L runs $U(\langle M \rangle, \langle M \rangle)$ while counting its steps up to $T(n)^{c_U}$. (Specifically, M_L first writes the binary representation of the integer $T(n)^{c_U}$ on the second tape, and initializes the third tape to the value zero. After simulating each step of U the machine M_L increments the counter on the third tape and compares it to the value on the second tape, stopping when the two values are equal.) At this point, if U outputted a bit b then M_L outputs $1 - b$, otherwise it outputs 0.

By definition, M_L decides L . To see that M_L runs in time $\text{poly}(T)$, we check that each of its steps runs in such time:

- Checking if the input is a valid encoding of a TM can be done in time $O(n \cdot \log(n))$,¹ which is at most $O(T(n) \cdot \log(T(n)))$ because $T(n) \geq n$.

¹Specifically, we first check whether the input parses correctly as Σ, S, δ , and then check that for every pair $s \in S, \sigma \in \Sigma$ there is a valid entry in the description of δ corresponding to input (s, σ) . Naively this can be done in cubic time, and by requiring that Σ, S, δ will be written in a suitably ordered way this can be done in time $O(n \cdot \log(n))$. (E.g., we first check that the input length is at least $|S| \cdot |\Sigma|$, otherwise we reject. Then we check that S and Σ are written in lexicographic order, and that the entries in the description of δ are written in lexicographic order of the inputs to δ , otherwise we reject. The ordering facilitates scanning over all pairs (s, σ) in linear time in the input length, for each entry we check that the output of δ (whose length is logarithmic in the input length) is of valid format.)

- Computing $T(n)^{c_U}$ can be done in time $O(T(n)) + \text{polylog}(T(n))$, because T is time-constructible and raising the integer $T(n)$ to the power c_U can be done in polynomial time in the integer's description length, which is $O(\log(T(n)))$.
- The machine simulates U for at most $T(n)^{c_U}$ steps. For each step, it increments a counter (whose value is at most $T(n)^{c_U}$) and compares it to the integer $T(n)^{c_U}$. Both the counter and the latter integer have description length at most $O(\log(T(n)))$, and incrementing and comparing can be done in linear time $O(\log(T(n)))$. Hence, simulating U takes time $O(T(n)^{c_U} \cdot \log(T(n)))$.

Each of the three steps above runs in time at most $\text{poly}(T(n))$, and the final step involves erasing the main worktape (whose content has at most $\text{poly}(T)$ bits), writing a bit and halting. So the overall running time is at most $\text{poly}(T(n))$. \blacksquare